

Arsenii Ashukha

[Home page](#) / [Google Scholar](#) / [GitHub](#)

e-mail: ars.ashuha@gmail.com tel: +79651782915

I'm a Research Scientist at Samsung AI Center Moscow. I (almost) received a PhD in Machine Learning. The results of my PhD were published at ICLR / ICML / NeurIPS and contributed to sparsification, uncertainty estimation, ensembling, and fundamentals of Bayesian deep learning (e.g., generative modeling priors, and incremental learning).

Prior to that, I was a part of Yandex Research in collaboration with the University of Amsterdam (with **Dmitry Vetrov** and **Max Welling**). I did ML engineering internships at Yandex (deep learning for music), Rambler (recommendation systems), and worked on NLP with **Natalia Loukachevitch**.

EDUCATION

- 2017-2021 PhD in ML, **National Research University Higher School of Economics**
 - PhD Thesis: Prior Knowledge for Deep Learning
 - Advisor: [Dmitry Vetrov](#)
- 2015-2017 MSc in Computer Science, **Moscow Institute of Physics and Technology (Cumlaude)**
- 2011-2015 BSc in Computer Science **Bauman Moscow State Technical University**

PROFESSIONAL EXPERIENCE

- *Research Scientist, Samsung AI Center* (2018 - Now):
I'm working on the development of deep learning algorithms. Specifically, my work is focused on uncertainty estimation, robustness, and fundamentals of Bayesian deep learning. I also contributed to computer vision research e.g., image inpainting.
- *Research Scientist, Yandex.Research & University of Amsterdam* (2016 - 2018):
Created Sparse Variational Dropout a method for sparsification of DNNs that, for the first time, allowed to achieve over 250x compression ratio (results published at ICML'17). The modified version of SparseVD with neuron-level sparsity allowed to accelerate inference of a CNN by 2-5 times and was involved in acceleration of the feature extraction for real image retrieval system (published at NeurIPS).
- *Machine Learning Engineer Intern, Yandex.Music* (summer of 2016):
I worked on feature extraction techniques for music data with convolutional neural networks. I also developed an evaluation of learned representations. The representations were used in the content-based recommendation system for yandex music.
- *Machine Learning Engineer Intern, Rambler&Co* (May - Oct 2015):
Worked on demographic classification and recommendation systems. My responsibility included improving the quality and performance of classifiers, automatic feature extraction algorithms, and recommendation algorithms. Stack of technologies: Hadoop, Hive, Spark, XGboost, VW, gensim.

SELECTED PUBLICATIONS

Google Scholar: scholar.google.com/citations?user=IU-kuP8AAAAJ

Most representative works are highlighted

- *Pitfalls of In-Domain Uncertainty Estimation and Ensembling in Deep Learning*, ICLR 2020 [[arXiv](#)]
[Arsenii Ashukha](#), Alexander Lyzhov, Dmitry Molchanov, Dmitry Vetrov
- *Variational Dropout Sparsifies Deep Neural Networks*, ICML 2017 [[arXiv](#)]
Dmitry Molchanov, [Arsenii Ashukha](#), Dmitry Vetrov
- *Greedy Policy Search: A Simple Baseline for Learnable Test-Time Augmentation*, UAI 2020 [[arXiv](#)]
Dmitry Molchanov, Alexander Lyzhov, Yuliya Molchanova, [Arsenii Ashukha](#), Dmitry Vetrov
- *The Deep Weight Prior*, ICLR 2019 [[arXiv](#)]
Andrei Atanov, [Arsenii Ashukha](#), Kirill Struminsky, Dmitry Vetrov, Max Welling
- *Variance Networks: When Expectation Does Not Meet Your Expectations*, ICLR 2019 [[arXiv](#)]
Kirill Neklyudov, Dmitry Molchanov, [Arsenii Ashukha](#), Dmitry Vetrov
- *Structured Bayesian Pruning via Log-Normal Multiplicative Noise*, NeurIPS 2017 [[arXiv](#)]
Kirill Neklyudov, Dmitry Molchanov, [Arsenii Ashukha](#), Dmitry Vetrov

CODE

- Check out very short and simple implementations of ML algorithms:
 - [Gradient Boosting](#) • [Real NVP normalizing flows](#) • [Quantile Regression DQN \(Distributional RL\)](#)
- Research-ready implementations:
 - [Multi-GPU SimCLRv1](#) closely reproduced results on both CIFAR-10 and ImageNet
 - [Ensembles \(Deep ensembles, Snapshot ensembles, cSGLD, FGE, etc.\)](#)

REVIEWING

- Conferences:
 - International Conference on Machine Learning, ICML (2019, 2020 top-33% highest-scored reviewers)
 - Neural Information Processing Systems, NeurIPS 2019 (top-50% highest-scored reviewers)
 - International Conference on Learning Representations, ICLR (2020, 2021)
- Workshops:
 - ICML Workshop on Invertible Neural Networks (2019, 2021, invertibleworkshop.github.io)
 - Bayesian Deep Learning Workshop (since 2017, bayesiandeeplearning.org)

TEACHING

- Supervisor of reading clubs on machine learning at HSE and Yandex school of data analysis (since 2017)
- A talk & a practical session on NFs at **Deep|Bayes** Summer School on Bayesian Deep Learning (since 2017)
- Lecturer, Moscow Institute of Physics and Technology: I was a lecturer and manager of the deep learning brunch of a faculty-wide machine learning course ~60 students (ml-mipt.github.io). Also, I taught deep learning and practical sessions on cutting-edge ML algorithms on a facultative course "Data Mining in

Action" ~ 200 students (<https://bit.ly/3eRLGYp>). The goal of this course is to make ML education available for everyone for free.

SUPERVISION

- Alexander Lyzhov (moved to NYU, PhD student), Deep Neural Network Ensembles: Analysis and Approaches to Diversification (MSc, 2020)
- Andrei Atanov (moved to EPFL, PhD student), Effective Learning of Deep Neural Networks Ensembles (BSc, 2018), Learning Deep Models with Small Data (MSc, 2020)
- Evgenii Nikishin (moved to Mila, PhD student), Stability Improvement and Knowledge Transfer in Deep Reinforcement Learning (MSc, 2019)

FRAMEWORKS & PROGRAMMING LANGUAGES

- I'm fluent in Python and I use to code in C/C++, Go, language is not a problem after all.
- I'm also fluent with common data science tools such as numpy, matplotlib, scikit-learn, numpy.
- I'm comfortable with common data science environment e.g., bash, git, linux.
- My primary deep learning framework is PyTorch. Prior to that, I had a decent experience with Theano+Lagange and TensorFlow.
- Comfortable with GPU clusters and distributed training.
- I have experience with MapReduce, Hadoop, Hive, and Spark.